

MAHENDRA ENGINEERING COLLEGE FOR WOMEN

ASSIGNMENT-2 SOLUTION

NAME OF THE SSTUDENT M.Gowsalya

REGISTER NUMBER:611419106027

YEAR/DEPARTMENT:1V-ECE

```
/data/user/0/com.microsoft.office.word/app_EmailAttachments2fd50a3fd
92b-4d15-b6e3-26f436f53ac3/DOC-20220926-WA0004..pdf#libraries import
pandas as pd import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
```

```
#load dataset
```

```
df = pd.read_csv(r"/content/Churn_Modelling.csv")
```

```
df.head(10)
```

```
  RowNumber  CustomerId  Surname  CreditLi
\
0           1    15634602  Hargrave
1           2    15647311    Hill
2           3    15619304    Onio
3           4    15701354    Boni
4           5    15737888  Mitchell
5           6    15574012    Chu
6           7    15592531  Bartlett
7           8    15656148    Obinna
8           9    15792365    He
9          10    15592389    H?
```

```

3    93826.63    0
4    79084.10    0
5    149756.71    1
6    10062.80    0
7    119346.88  1  8  74940.50  0
9           71725.73    0

```

```
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   RowNumber             10000 non-null  int64
 1   CustomerId            10000 non-null  int64
 2   Surname               10000 non-null  object
 3   CreditScore           10000 non-null  int64
 4   Geography             10000 non-null  object
 5   Gender                10000 non-null  object
 6   Age                   10000 non-null  int64
 7   Tenure                10000 non-null  int64
 8   Balance               10000 non-null  float64
 9   NumOfProducts         10000 non-null  int64
10   HasCrCard             10000 non-null  int64
11   IsActiveMember        10000 non-null  int64
12   EstimatedSalary       10000 non-null  float64
13   Exited                10000 non-null  int64
dtypes: float64(2), int64(9), object(3)
memory usage: 1.1+ MB


```

```


#Visualizations #Univariate
Analysis import seaborn as
sns
sns.kdeplot(df['CreditScore'])

```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7fc4a0cd2790>
```




```
#Bi - Variate Analysis  
plt.bar(df.CustomerId, df.CreditScore)  
plt.title('CreditScore')  
plt.xlabel('CustomerId')  
plt.ylabel('CreditScore')  
  
Text(0, 0.5, 'CreditScore')
```




```
sns.lmplot(x='Tenure', y='Balance', data=df ,hue='Exited',size=8)
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/regression.py:581:  
UserWarning: The `size` parameter has been renamed to `height`;  
please update your code. warnings.warn(msg, UserWarning)
```

```
<seaborn.axisgrid.FacetGrid at 0x7fc4a149e2d0>
```



```
#Multi - Variate Analysis ax =
df[["CreditScore","Age","Tenure","Balance"]].plot(figsize=(80,40)
) ax.legend(loc='center left', bbox_to_anchor=(1, 0.5));
```




```
df.isnull().sum()
```

RowNumber	0
CustomerId	0
Surname	0
CreditScore	0
Geography	0
Gender	0
Age	0
Tenure	0
Balance	0
NumOfProducts	0
HasCrCard	0
IsActiveMember	0
EstimatedSalary	0
Exited	0

```
dtype: int64
```

```
plt.figure(figsize=(15,13))  
sns.heatmap(df.corr(),annot=True,cmap='BuPu')  
plt.show()
```



```
df.drop(['RowNumber', 'CustomerId', 'Surname'], axis=1, inplace=True)
```

```
df.head()
```

```

    CreditScore Geography Gender Age Tenure      Balance NumOfProducts
\
0          619    France Female   42      2         0.00
1
1          608     Spain Female   41      1  83807.86
1
2          502    France Female   42      8 159660.80
3
3          699    France Female   39      1         0.00
2

```

```

4           850           Spain Female           43           2 125510.82
1

```

```

      HasCrCard IsActiveMember EstimatedSalary Exited 0 1
1           1      101348.88      1
1           0           1      112542.58                      0
2           1           0      113931.57 3           0           0      1
      93826.63                      0
4           1                      1      79084.10      0

```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 10000 entries, 0 to 9999
```

```
Data columns (total 11 columns):
```

```
# Column Non-Null Count Dtype
```

```

-----
0  CreditScore      10000 non-null int64
1  Geography      10000 non-null object
2  Gender          10000 non-null object
3  Age            10000 non-null int64
4  Tenure          10000 non-null int64
5  Balance         10000 non-null float64
6  NumOfProducts   10000 non-null int64
7  HasCrCard       10000 non-null int64
8  IsActiveMember  10000 non-null int64
9  EstimatedSalary 10000 non-null float64
10 Exited          10000 non-null int64

```

```
dtypes: float64(2), int64(7), object(2)
```

```
memory usage: 859.5+ KB
```

```
df["Geography"].unique() array(['France', 'Spain',
```

```
'Germany'], dtype=object) df["Gender"].unique()
```

```
array(['Female', 'Male'], dtype=object)
```

```
geo=pd.get_dummies(df["Geography"],drop_first=False
```

```
) geo.head()
```

```

      France Germany Spain
0           1           0           0
1           0           0           1
2           1           0           0
3           1           0           0

```



```
4          0          0          1
```

```
gen=pd.get_dummies(df["Gender"],drop_first=False  
) df=pd.concat([df, geo,gen], axis=1) df
```

```
      CreditScore Geography Gender Age Tenure      Balance  
NumOfProducts \
```

0

	619	France Female 42	2	0.00
1				
1	608	Spain Female 41 1	83807.86	
1				
2	502	France Female 42	8 159660.80	
3				
3	699	France Female 39	1 0.00	
2				
4	850	Spain Female 43	2 125510.82	
1				
...
...				
9995	771	France Male 39	5 0.00	
2				
9996	516	France Male 35	10 57369.61	
1				
9997	709	France Female 36	7 0.00	
1				
9998	772	Germany Male 42	3 75075.31	
2				
9999	792	France Female 28	4 130142.79	
1				

	HasCrCard	IsActiveMember	EstimatedSalary	Exited	France
Germany \					
0	1	1	101		
0					
1	0	1	112		
0					
2	1	0	113		
0					
3	0	0	93		

1

0

4 1 1 79

0

...

...

9995 1 0 96

0

9996 1 1 101

0

1 1 0

2 0 1 0

3 0 1 0

4 1 1 0

...

9995 0 0 1

9996 0 0 1

9997 0 1 0

9998 0 0 1

9999 0 1 0

[10000 rows x 16 columns]

df.drop(["Geography","Gender"], axis=1, inplace=True)

df.head()

	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	\
0	619	42	2	0.00	1	1	
1	608	41	1	83807.86	1	0	
2	502	42	8	159660.80	3	1	
3	699	39	1	0.00	2	0	
4	850	43	2	125510.82	1	1	

	IsActiveMember	EstimatedSalary	Exited	France	Germany	Spain	Female	\
0	1	101348.88	1	1	0	0		

2

1

1 1 112542.58 0 0 0 1

1

2 0 113931.57 1 1 0 0

1

3 0 93826.63 0 1 0 0

1

4 1 79084.10 0 0 0 1

1

Male

0 6

1 0

2 0

3 3

4 3

x=df.drop('Exited',axis=1)

x

CreditScore Age Tenure Balance 0 619

42 2 0.00 502 42 8 159660.80

3 1

3 699 39 1 0.00 2 0

4 850 43 2 125510.82 1 1

.... .

9995 771 39 5 0.00 2 1

9996 516 35 10 57369.61 1 1

9997 709 36 7 0.00 1 0

9998 772 42 3 75075.31 2 1

9999 792 28 4 130142.79 1 1

IsActiveMember EstimatedSalary France Germany Spain Female

Male

0 1 101348.88 1 0 0 1

0

1 1 112542.58 0 0 1 1

0

2 0 113931.57 1 0 0 1

0

3

3	0	93826.63	1	0	0	1
0						
4	1	79084.10	0	0	1	1
0						
...
...						
9995	0	96270.64	1	0	0	0
1						
9996	1	101699.77	1	0	0	0
1						
9997	1	42085.58	1	0	0	1
0						
9998	0	92888.52	0	1	0	0
1						
9999	0	38190.78	1	0	0	1
0						

[10000 rows x 13 columns]

y=df['Exited']

y

0	1
1	0
2	1
3	0
4	0
	..
9995	0
9996	0
9997	1
9998	1

```

9999      0
Name: Exited, Length: 10000, dtype: int64

df.shape (10000, 14)

x.shape
(10000, 13)

y.shape (10000,) from sklearn.model_selection
import train_test_split
x_train,x_test, y_train,y_test =
train_test_split(x,y, test_size=0.2,random_state=0)
x_train.shape (8000, 13) x_test.shape (2000, 13)
y_test.shape (2000,) from sklearn.preprocessing
import StandardScaler sc = StandardScaler() x_train =
sc.fit_transform(x_train) x_train
array([[ 0.16958176, -0.46460796,  0.00666099, ...,  1.74309049,
         1.09168714, -1.09168714],
       [-2.30455945,  0.30102557, -1.37744033, ..., -0.57369368,
        -0.91601335,  0.91601335],
       [-1.19119591, -0.94312892, -1.031415 , ..., -0.57369368,
         1.09168714, -1.09168714],
       ...,
       [ 0.9015152 , -0.36890377,  0.00666099, ..., -0.57369368,
        -0.91601335,  0.91601335],
       [-0.62420521, -0.08179119,  1.39076231, ...,  1.74309049,
         1.09168714, -1.09168714],
       [-0.28401079,  0.87525072, -1.37744033, ..., -0.57369368,
         1.09168714, -1.09168714]])

x_test = sc.transform(x_test)
x_test
array([[ -0.55204276, -0.36890377,  1.04473698, ..., -0.57369368,
         1.09168714, -1.09168714],
       [-1.31490297,  0.10961719, -1.031415 , ..., -0.57369368,
         1.09168714, -1.09168714],
       [ 0.57162971,  0.30102557,  1.04473698, ...,  1.74309049,
         1.09168714, -1.09168714],
       ...,
       [-0.74791227, -0.27319958, -1.37744033, ...,  1.74309049,
        -0.91601335,  0.91601335],

```

```
[-0.00566991, -0.46460796, -0.33936434, ..., -0.57369368,  
 -0.91601335, 0.91601335],  
[-0.79945688, -0.84742473, 1.04473698, ..., -0.57369368,  
 -0.91601335, 0.91601335]])
```